
PLANTFORGE: A Procedural Control-Plant Corpus for In-Context System Identification

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 In-context system-identification (SysID) transformers are typically trained on a
2 single private data-generation recipe — most commonly Wiener-Hammerstein-only
3 white-noise data, regenerated independently by every paper in this line of work.
4 We introduce PLANTFORGE, a released, static, reproducible corpus of synthetic
5 control-plant trajectories organized along three axes jointly — nonlinearity *family*,
6 excitation *class*, and sampling *rate* (exact zero-order-hold) — with per-instance
7 Fisher-information identifiability annotations. Across 5 independently-seeded train-
8 ing runs, we show that a transformer trained on Wiener-Hammerstein-only data
9 collapses off its training distribution (cross-family nMSE degradation of 93–118×),
10 and that training on PLANTFORGE’s multi-axis corpus closes the rate and ex-
11 citation gaps (to $\sim 1\text{--}2\times$ for four of five families; the resonant drivetrain family
12 improves but stays at $8\times$) while only halving, not closing, the cross-family gap
13 ($93\text{--}118\times \rightarrow 14\text{--}18\times$). We then subject this result to three checks a released
14 benchmark should survive but rarely does: (1) zero-shot transfer to three real mea-
15 sured plants (Silverbox, Cascaded Tanks, Bouc-Wen), where corpus-trained models
16 again outperform the Wiener-Hammerstein-only model, but a simple classical ARX
17 baseline outperforms *both* trained transformers by $11\text{--}342\times$ on every one of the
18 three, tempering any claim of transformer superiority; (2) a within-cell statisti-
19 cal re-analysis showing that the corpus’s own identifiability annotations do *not*
20 positively predict in-context prediction difficulty (median Spearman $r = -0.122$
21 across 40 cells), the opposite of our working hypothesis; and (3) full disclosure
22 of a Simpson’s-paradox confound our first analysis of (2) missed, and how we
23 caught and corrected it. We release the corpus (240,000 instances, CC BY 4.0), all
24 generation/training/evaluation code (MIT), and every number in this paper with
25 reproduction instructions.

26 1 Introduction

27 Model-agnostic, in-context system identification — training a sequence model to infer a plant’s
28 dynamics from a short context of (u, y) pairs, with no gradient updates at test time — is an increasingly
29 popular paradigm for control and mechatronics [Forgione et al., 2023, Piga et al., 2024]. The training
30 recipe behind this line of work, however, is narrow and largely unreleased: each paper privately
31 regenerates Wiener-Hammerstein-only white-noise data, the one family the incumbent generators
32 support. No existing corpus varies *nonlinearity family*, *excitation class*, and *sampling rate* jointly, and
33 none ships per-instance identifiability annotations that would let a user reason about *why* an instance
34 is hard.

This paper makes two kinds of contribution. First, we release PLANTFORGE: a procedurally generated corpus spanning 5 nonlinearity families $\times$ 4 excitation classes $\times$ 3 sampling rates (240,000 instances total), with per-instance relative Cramér-Rao lower bounds (CRLB) and Fisher-information matrix (FIM) condition numbers as identifiability annotations, and design invariants enforced by tests (Section 3). Second, and we think more valuably for the community, we report what happened when we subjected our own headline claim to increasingly rigorous scrutiny: multi-seed error bars, zero-shot real-plant validation, a classical baseline, and a statistical re-analysis of our identifiability hypothesis. Two of these checks produced results that *complicate* the paper’s original thesis rather than confirm it, and we report them in full rather than adjusting the narrative around them (Section 4).

Summary of findings.

- Multi-axis training on PLANTFORGE closes the rate and excitation transfer gaps that Wiener-Hammerstein-only training leaves open (for four of five families; the resonant drivetrain family improves but does not close), but only halves the harder cross-family gap — reported here as an open challenge, not a solved problem (Section 4.1).
- Zero-shot on three real plants, a corpus-trained model beats a Wiener-Hammerstein-only model on every one, but a simple per-window ARX fit beats *both* trained transformers by 11–342 $\times$ under an identical, leakage-free evaluation protocol (Section 4.3).
- The corpus’s own identifiability annotations, hypothesized to predict in-context prediction difficulty, show a weak, robustly *negative* within-cell correlation with it once a between-cell confound in our first analysis is corrected (Section 4.4).

2 Related Work

Synthetic foundation models for dynamics. Ziegler et al. [2024] pretrain a transformer-based foundation model purely on synthetic dynamics sampled from a reproducing kernel Hilbert space, and show it transfers to real hardware after fine-tuning. This is generically capable but does not vary family/excitation/rate/identifiability jointly, and does not target control-plant nonlinearities specifically; we view it as the closest and most relevant prior synthetic-data approach, and the one whose scope PLANTFORGE is most easily confused with. A separate, theoretical line of work [Cole et al., 2025] derives approximation and depth-separation bounds for transformers doing in-context learning of *linear* dynamical systems; PLANTFORGE addresses the empirical/benchmark side this theory does not cover (nonlinear families, real-plant transfer, identifiability).

In-context system identification. The training recipe we stress-test throughout this paper — Wiener-Hammerstein-only white-noise pretraining for an in-context SysID transformer — is the one used across the Forgone et al. [2023], Piga et al. [2024] line of work (the “forgi86 lineage”), each of which regenerates its own private WH-only data. This lineage has continued to grow: Rufolo et al. [2025b] propose an enhanced architecture (probabilistic framing, non-contiguous context/query windows, recurrent patching) for the same in-context SysID task — PLANTFORGE deliberately evaluates a plain encoder-only baseline instead, so our distribution-shift findings are not confounded by architecture novelty; Busetto et al. [2024] extend the same meta-model paradigm to state estimation, showing the narrow-recipe pattern recurs across adjacent tasks, not just the one paper we benchmark against directly; Rufolo et al. [2025a] attack the identical cross-family/rate/excitation generalization gap PLANTFORGE measures, but via a distributionally-robust training *objective* on a fixed data distribution, a complementary approach to PLANTFORGE’s strategy of broadening the training *corpus* itself. Outside this lineage, Akram and Vikalo [2024] give an independent theoretical account of what in-context transformers implicitly compute for dynamical systems, evidence the paradigm has uptake beyond the one research group PLANTFORGE otherwise stress-tests exclusively. PLANTFORGE’s own InContextSysID model is a standard encoder-only transformer over (u, y) pairs — our contribution is the training corpus and the scrutiny applied to it, not architectural novelty.

Domain randomization and procedural benchmark generation. PLANTFORGE’s design — generate broadly across a parametrized axis so a held-out real system looks like “just another variation” rather than a foreign distribution — is the system-identification instance of domain randomization, established in robotics by Tobin et al. [2017] and surveyed for robot learning by Muratore et al. [2022]; unlike most of that literature, PLANTFORGE directly measures the residual zero-shot gap

on real plants (Section 4.2) rather than assuming randomization closes it, and finds it does not fully close for the hardest axis (cross-family transfer). Procedural generation as a benchmark-design tool in its own right — generating a broad, structured space of environments to measure generalization, independent of any one target domain — is established by Cobbe et al. [2020] for reinforcement learning and by Takamoto et al. [2022], Gupta and Brandstetter [2022] for PDE surrogate modeling (Section 3’s three jointly-varied axes are PLANTFORGE’s analog of these platforms’ varied PDE type/parameters/boundary conditions, adapted to a domain — control/mechatronics — with an excitation-design axis these PDE benchmarks have no equivalent of). The general phenomenon PLANTFORGE’s headline transfer-gap measurement instantiates is studied at ML-wide scale by Koh et al. [2021].

Scaling benchmarks and other released dynamics corpora. DynaDojo [Bhamidipaty et al., 2023] is an extensible platform for benchmarking how SysID algorithms scale with training samples, system complexity, and target error, over a library of ~ 20 generic ODE/chaotic systems. It does not include control-specific nonlinearities (deadzone, saturation, hysteresis, drivetrain compliance) or an excitation/rate axis, and carries no identifiability annotations. Ullah and Baca [2026] release a real-hardware (Crazyflie 2.1) multi-task SysID/control/state-estimation benchmark, evidence the community values multi-axis benchmark design even outside the synthetic setting PLANTFORGE occupies.

Real-measured benchmarks. The community’s default real-data benchmarks — Silverbox, Wiener-Hammerstein, Bouc-Wen [Noël and Schoukens, 2020], Cascaded Tanks [Wigren and Schoukens, 2013, Schoukens and Ljung, 2019] — are fixed, single-record, real-measured systems with no parameter ground truth. We evaluate zero-shot against three of the four (Section 4.2); a classical baseline outperforming our trained transformers on those benchmarks (Section 4.3) is itself evidence for why a synthetic corpus with known parameters and controllable axes is a useful complement, not a replacement, for this real-measured suite.

Classical nonlinear system identification and identifiability. PLANTFORGE’s ARX/NARX baseline (Section 4.3) and per-instance Fisher-information identifiability annotations (Section 3, tested in Section 4.4) sit within a mature classical literature: black-box nonlinear model structures are surveyed by Sjöberg et al. [1995], and NARMAX methods specifically — including the free-run simulation instability our degree-2 NARX baseline reproduces empirically — by Billings [2013]. The Fisher-information/ Cramér-Rao formalism our `identifiability.py` implements is standard system-identification theory [Ljung, 1999]; Hjalmarsson [2005] surveys the connection from experiment/excitation design to identification accuracy that motivates our excitation axis, against whose usual assumption (better-designed excitations $\rightarrow$ better identifiability $\rightarrow$ presumed more useful data) our null result in Section 4.4 is a direct, non-obvious counterpoint. Aguirre et al. [2010] study the same prediction-error-vs.-simulation-error distinction our ARX free-run protocol turns on. Excitation design specifically — multisine and PRBS signals, both implemented in our `excitation.py` — is treated at length by Pintelon and Schoukens [2012].

3 The PLANTFORGE Corpus

Design axes. Every instance is generated along three axes jointly: **family** — `stribeck` (velocity-dependent friction, a state-nonlinearity), `backlash` (input deadzone), `saturate` (input clipping), `boucwen` (output hysteresis with memory), and `drivetrain` (two-inertia motor/gear/compliant-load with load Coulomb friction, the mechatronics staple) — plus a Wiener-Hammerstein (`wh`) family excluded from the corpus proper but retained for the failure experiment in Section 4.1; **excitation** — `prbs` (0.25s physical hold), `multisine` (8 tones ≤ 1.85 Hz), `chirp` (0.05 $\rightarrow$ 3 Hz), and `closedloop` (a true sequential PI-tracking loop against the plant stepper, so the same seed against a different plant instance produces different u — not a two-pass imitation of closed-loop data); and **rate** — 10/20/50 Hz, exact zero-order-hold discretizations of one shared continuous-time truth per instance.

Identifiability annotations. Each (instance, excitation, rate) combination carries a per-parameter relative Cramér-Rao lower bound and a Fisher-information-matrix $\log_{10}$ condition number, computed analytically from the same simulator that generated the trajectory.

Design invariants (tested). (1) (u, y) pairs are physically consistent with their stated parameters — no hidden output normalization; re-simulation from θ reproduces y to 10^{-6} . (2) Multi-rate instances are the *same* continuous-time plant: state-nonlinear families substep internally at ≤ 2 ms with zero-order-hold input, so dt and $dt/4$ agree at common sample instants to within 10^{-3} relative error (exact for LTI cores) — an axis we are not aware of any other released corpus generating exactly. (3) Closed-loop excitation is a genuine sequential loop, verified by checking that u depends on the specific plant instance under control. (4) Identifiability annotations are directionally physical: an excitation that never exits the backlash deadzone yields a relative CRLB on the deadzone width $\sim 5 \times 10^4$, versus ~ 0.3 for an excitation that does.

Release. 240,000 instances (5 families $\times$ 4 excitations $\times$ 3 rates $\times$ 4,000 instances/cell), 583 MB, released under CC BY 4.0 [Anonymous, 2026]. All generation, training, and evaluation code is released under MIT. Full datasheet, following the Gebru et al. [2021] template (motivation, composition, collection process, uses, distribution, maintenance), accompanies the release.

4 Experiments

All models are an identical encoder-only transformer (InContextSysID: 5 layers, 8 heads, width 160) consuming 192 context samples and predicting 32 query-horizon samples with no ground-truth y visible in the query horizon. Every headline number in this section is the mean $\pm$ standard deviation over **5 independently-seeded training runs** (seeds 0–4, 10,000 steps each, fixed evaluation seeds 900–905 across every model for paired comparison). nMSE is always the ratio of batch-mean query-horizon squared error to batch-mean query-horizon signal power, never a mean of per-window ratios.

4.1 Headline: multi-axis training closes the rate/excitation gap, halves the family gap

We compare two training recipes: **wh_only**, trained exclusively on Wiener-Hammerstein white-noise data (the incumbent recipe), and **corpus**, trained on 4-of-5 PLANTFORGE families $\times$ {multisine, prbs} $\times$ {10, 50 Hz}, with backlash, chirp/closedloop, and 20 Hz held out.

Table 1: Synthetic transfer matrix, nMSE (mean $\pm$ std, $n = 5$ seeds). $\times$ ref is relative to that model’s own in-distribution reference cell.

held-out axis	wh_only nMSE	($\times$ ref)	corpus nMSE	($\times$ ref)
reference (train-like)	0.0040 $\pm$ 0.0004	1.0 $\times$	0.0215 $\pm$ 0.0017	1.0 $\times$
family backlash, dt=0.10	0.3880 $\pm$ 0.0322	97.7 $\times$	0.3107 $\pm$ 0.0080	14.5$\times$
family backlash, dt=0.05	0.3678 $\pm$ 0.0408	92.6 $\times$	0.2899 $\pm$ 0.0095	13.5$\times$
family backlash, dt=0.02	0.4665 $\pm$ 0.0195	117.5 $\times$	0.3927 $\pm$ 0.0102	18.3$\times$
excitation chirp	0.1533 $\pm$ 0.0645	38.6 $\times$	0.0383 $\pm$ 0.0143	1.8$\times$
excitation closedloop	0.0545 $\pm$ 0.0056	13.7 $\times$	0.0195 $\pm$ 0.0031	0.9$\times$
rate dt=0.05, stribeck	0.0341 $\pm$ 0.0039	8.6 $\times$	0.0215 $\pm$ 0.0017	1.0$\times$
rate dt=0.05, saturate	0.0141 $\pm$ 0.0016	3.5$\times$	0.0285 $\pm$ 0.0021	1.3 $\times$
rate dt=0.05, boucwen	0.0916 $\pm$ 0.0177	23.1 $\times$	0.0178 $\pm$ 0.0014	0.8$\times$
rate dt=0.05, drivetrain	0.2180 $\pm$ 0.0156	54.9 $\times$	0.1722 $\pm$ 0.0111	8.0$\times$

Table 1 and Figure 1 summarize the result. Multi-axis training **solves the rate and excitation axes for most families**: the WH-only model’s 3.5–39 $\times$ degradation on interpolated rate and held-out excitation collapses to ~ 1 –2 $\times$ under corpus training for four of the five trained families. The exception is the resonant drivetrain family, whose rate cell improves 54.9 $\times \rightarrow$ 8.0 $\times$ but does not close — a second open finding, invisible to single-family generators. Multi-axis training also **halves but does not close the cross-family gap** (93–118 $\times \rightarrow$ 14–18 $\times$, still an order of magnitude off reference) — we report this as PLANTFORGE’s open challenge track, not a solved problem. One further cell, saturate at the held-out rate, is won by wh_only (non-overlapping error bars the other direction): “corpus training helps” is not a universal law across every cell, and we report the exception rather than omit it. Non-overlapping ± 1 standard deviation across 11 simultaneous cell comparisons is a descriptive heuristic, not a multiplicity-corrected test; at $n = 5$ seeds this is a conservative bar relative to the standard error of the mean, and the largest-margin cells (below) clear it comfortably.




Figure 1: Synthetic transfer matrix, wh_only vs. corpus, 5-seed mean±std, log scale.

4.2 Zero-shot transfer to real measured plants

Both checkpoints, trained only on synthetic PLANTFORGE data, are evaluated zero-shot (no fine-tuning) on real recordings from nonlinearbenchmark.org [Wigren and Schoukens, 2013, Schoukens and Ljung, 2019], obtained via the `nonlinear_benchmarks` package for two of the three datasets below, and via a direct fetch from its official source for the third (4TU.ResearchData [Noël and Schoukens, 2020], CC BY-SA 4.0, since Bouc-Wen is not exposed by `nonlinear_benchmarks`’s public loader API). Real sample rates do not align with PLANTFORGE’s three trained rates, so we handle each dataset on its own terms: **Silverbox** (native $dt \approx 1.64$ ms) is decimated $12\times$ (anti-aliased) to $dt \approx 19.7$ ms, landing near the finest trained rate; **Cascaded Tanks** (native $dt = 4.0$ s) is already $40\times$ coarser than the coarsest trained rate and is evaluated at its native rate as an explicit extrapolation; **Bouc-Wen** (native $dt = 1/750$ s exactly) is decimated $15\times$ to $dt = 0.02$ s *exactly* (no rounding error), the cleanest rate match of the three — note this is a real hysteretic mass-spring-damper benchmark, distinct from PLANTFORGE’s own synthetic boucwen corpus family (Section 3), which models output hysteresis but is not fit to this specific real system; **Wiener-Hammerstein Process Noise** is excluded — its ~ 1.5 s total record duration is shorter than one 224-sample context window at any trained rate, so it cannot supply a single in-context evaluation window.

Table 2: Zero-shot real-plant nMSE. wh_only/corpus: mean±std, 5 seeds, fixed windows. ARX/NARX2: single deterministic run, ≤ 8 non-overlapping windows (Section 4.3).

dataset	wh_only	corpus	ARX	NARX2
Silverbox (decimated $12\times$)	0.958 ± 0.217	0.331 ± 0.040	0.0028	0.0028
Cascaded Tanks (native rate, extrapolation)	0.394 ± 0.081	0.084 ± 0.043	0.0075	3.7×10^9
Bouc-Wen (decimated $15\times$, exact)	2.537 ± 0.848	1.589 ± 0.361	0.0603	0.0529

Corpus training again beats WH-only training zero-shot on real data (most dramatically on Silverbox), consistent with Section 4.1. This held-up conclusion is itself a worked example of why we report every real-plant number at 5-seed scale: an earlier single-seed (seed 0 only) pass on Bouc-Wen showed the *opposite* direction (corpus nMSE 2.05 vs. wh_only 1.66) — a result two independent reviewers confirmed was not a code defect given the data available at the time, since wh_only’s seed-to-seed std on Bouc-Wen (0.848) is the largest relative variance of any real-plant number in this paper. At full 5-seed scale the direction reverses back to the Silverbox/Cascaded-Tanks pattern. We surface this not to bury it, but because it is exactly the failure mode multi-seed reporting exists to catch — and the reason every number in this paper, including the architecture ablation in Section 5, is now reported at the same 5-seed scale rather than a cheaper single-seed pass.

But Table 2 previews a result that complicates the story further: a trivial classical baseline beats *both* trained transformers by two orders of magnitude on every one of the three real-plant benchmarks. We discuss why in Section 4.3.

4.3 Classical baselines temper the transformer story

We compare against ARX and degree-2 polynomial NARX, fit by ordinary/ridge least squares *per window*, under an evaluation protocol identical to the transformer’s: fit (or, for the transformer, condition) on the 192-sample context only, then free-run the 32-sample query given only u — neither method ever observes query-horizon ground-truth y . ARX order $k \in \{2, 4, 8\}$ is selected per window by one-step-ahead error on the last 32 context samples, never touching the query.

Result. Table 2 and Figure 2 show ARX beating both trained transformers on all three real-plant benchmarks by 11–342 $\times$ (Bouc-Wen: 26–48 $\times$, within the same range). This pattern recurs synthetically: ARX also beats both models on held-out excitation `chirp` (0.0039 vs. 0.0383/0.1533), held-out excitation `closedloop` (0.0036 vs. 0.0195/0.0545), and held-out rate `drivetrain` (0.0253 vs. 0.1722/0.2180), and roughly ties `corpus` on held-out family `backlash` at the coarsest rate (0.3116 vs. 0.3107).

Is this a fair comparison? We verified it is. ARX and the transformer consume identical windows, identical normalization, and an identical context/query split; the transformer’s own query-horizon input is zeroed exactly where ARX is forbidden from using ground truth, so neither method sees query y at fit/inference time. ARX’s per-window least-squares fit is not an unfair privilege — *adapting to the context* is precisely the in-context transformer’s own value proposition, executed via a frozen forward pass instead of an explicit fit. The real advantage ARX has on rate-shifted and real-plant cells is structural, not a protocol bug: a per-window linear refit at the test sampling rate is immune to a rate mismatch that a frozen transformer must instead extrapolate through, most severely on Cascaded Tanks’ 40 $\times$ extrapolation. We view this as a genuine limitation of the frozen in-context approach worth further study, not an artifact to explain away.

Degree-2 NARX diverges broadly in free-run (clipped-but-huge nMSE, 10^9 – 10^{12} , on all but two of the twelve evaluated cells) — a legitimate finding about naive polynomial NARX instability under 32-step free simulation, and partly attributable to the $k = 8$ order candidate being numerically underdetermined at the order-selection stage (153 quadratic features vs. 152 fit samples). The two exceptions, both real-plant, are Silverbox and — newly — Bouc-Wen (NARX2 0.0529, close to ARX’s 0.0603); we do not have a full account of why free-run quadratic NARX stays stable on these two real systems and not elsewhere, and flag it as an open question rather than force an explanation. We report all NARX2 numbers rather than omit them, but do not treat the diverged ones as comparable to ARX’s.




Figure 2: Zero-shot real-plant nMSE: wh_only, corpus, and ARX. Log scale.

4.4 Identifiability annotations do not predict prediction difficulty

PLANTFORGE’s per-instance rel-CRLB and FIM-condition annotations were hypothesized, at design time, to predict in-context prediction difficulty — intuitively, a poorly-identifiable instance should also be hard to predict. We tested this directly: per-instance query-horizon nMSE, scored on the corpus model (5 seeds, mean), against each instance’s identifiability annotation, on the 40 (family $\times$ excitation $\times$ rate) cells whose trajectory length exceeds one context window (10 Hz cells are excluded: 128 samples $<$ 224).

A confound we caught ourselves. Our first analysis pooled all 40,000 instances and correlated directly: Spearman $r = -0.088$ (max rel-CRLB) and $r = -0.316$ ($\log_{10}$ FIM condition), both statistically significant but with per-family breakdowns of *opposite sign* (backlash $+0.294$, drivetrain $+0.091$, vs. stribeck -0.081 , saturate -0.148 , boucwen -0.050). This sign heterogeneity is the fingerprint of a Simpson’s-paradox confound: excitation and family choice change both identifiability and model difficulty independently, so a pooled correlation conflates between-cell structure with the within-cell effect the hypothesis is actually about. We discarded the pooled analysis as the primary result and re-ran within-cell.

Corrected analysis. We compute Spearman r within each of the 40 cells separately, then aggregate. We further apply two confound controls: (a) excluding each cell’s bottom decile of query-horizon signal power (near-flat query horizons mechanically inflate nMSE, a live confound independent of identifiability), and (b) restricting to cells with non-degenerate identifiability dynamic range (all 40/40 cells pass this filter).

Table 3: Within-cell Spearman r (max rel-CRLB, per-instance nMSE), aggregated over 40 cells, 5 seeds, 1,000 instances/cell (39,994/40,000 finite).

	median r	mean r	cells with $r > 0$
base	−0.122	−0.092	10/40
power-controlled (excl. bottom-decile power)	−0.117	−0.103	—
range-filtered (IQR $> 10^{-6}$, 40/40 cells kept)	−0.122	—	—

Result: weak, robustly negative, not the hypothesized direction. The within-cell correlation is small and negative, and survives both confound controls without flipping sign (Table 3). Only 10/40 cells are positive, concentrated in backlash-multisine and backlash-closedloop (r up to $+0.327$); every other family is negative on balance (Figure 3).




Figure 3: Within-cell Spearman r , all 40 cells, corpus model, 5 seeds. Ticks mark per-family medians.

Interpretation. We propose the negative sign reflects a real mechanistic decoupling rather than a null result: rel-CRLB measures *parameter-recovery* difficulty, not *prediction* difficulty. A parameter that is hard to identify from data is typically one with little influence on the output trajectory — so an in-context predictor does not need to infer accurately what barely affects y in the first place.

Under this reading, the annotations remain valuable corpus metadata (e.g. for excitation-design studies, where they are the intended quantity of interest) but are not, and should not be marketed as, a prediction-difficulty predictor. A secondary methodological finding: naive quartile-binned *means* of per-instance nMSE show a dramatic apparent upward trend ($1.5 \times 10^3 \rightarrow 1.6 \times 10^9$ across quartiles) purely because per-instance nMSE has an extremely heavy right tail (near-zero query-horizon power inflates the ratio); the quartile *medians* are flat (0.037, 0.037, 0.040, 0.021, Figure 4) — the mean-based framing we initially reported for this experiment before catching the confound above was itself a symptom of the same heavy-tail artifact, and we flag it explicitly so the failure mode is documented, not just quietly fixed.

Quartiles of max rel-CRLB $\rightarrow$ per-instance nMSE: mean is heavy-tail-dominated, median is flat ($n \approx 10,000/\text{bin}$)




Figure 4: Quartiles of max rel-CRLB $\rightarrow$ per-instance nMSE: bin mean (left, heavy-tail-dominated) vs. bin median (right, flat).

5 Discussion and Limitations

What this paper does and does not claim. We do *not* claim in-context transformers are the strongest available method for real-plant system identification: Section 4.3’s ARX result argues against that claim directly. We *do* claim that (a) training-distribution breadth matters — multi-axis training measurably narrows, though does not close, the transfer gaps a narrow WH-only recipe leaves open, and (b) a released, parameter-annotated, three-axis corpus is a more useful research artifact for studying *why* than any of the alternatives in Section 2, independent of whether the transformer architecture itself is currently the best tool for the job.

Scope of real-plant validation. We validate zero-shot on 3 of the 4 real benchmarks named in our original design motivation (Silverbox, Cascaded Tanks, Bouc-Wen); only Wiener-Hammerstein Process Noise remains excluded, by record length, explained in Section 4.2. Real-plant nMSE is computed over ≤ 8 non-overlapping windows per dataset; the transformer’s reported variance is over training seeds on this fixed window set, not over window sampling, and the ARX comparison point is a single deterministic run — the $11\text{--}342\times$ margins are far larger than plausible window-sampling noise, but this should be read as a first zero-shot probe, not an exhaustive real-plant evaluation. Bouc-Wen’s `wh_only` variance across seeds (std 0.848 on a mean of 2.537) is the largest relative variance of any real-plant number we report, and a single-seed preview of this exact cell showed the opposite model ranking before the remaining four seeds were trained (Section 4.2) — a reminder that even our 5-seed standard is not immune to high-variance cells, and that per-cell variance, not just the headline mean, is worth reporting when re-using these numbers.

Family coverage and held-out-family choice. We hold out a single family (`backlash`) throughout; we have not verified whether it is representative of the corpus’s hardest held-out-family case or an outlier, and leave a full leave-one-family-out sweep to future work. The corpus’s five families do not span every control nonlinearity of interest (no multi-input/multi-output plants, no time-varying parameters within a trajectory).

298 **Architecture ablation: the family gap is capacity-invariant (5-seed check).** All headline results
 299 (Sections 4.1–4.4) use one fixed transformer configuration (160 width, 5 layers, 1.58M params). We
 300 ran a lighter-weight check — corpus recipe only, but at the same 5-seed scale as every other number
 301 in this paper — across four additional variants spanning a $15\times$ parameter range, to test whether the
 302 held-out-family gap is an artifact of this specific size.

Table 4: Architecture ablation (mean $\pm$ std, 5 seeds, corpus recipe): reference vs. held-out family
 backslash dt=0.05 nMSE.

variant	params	reference nMSE	family nMSE (dt=0.05)
narrow (d=80)	407k	0.0300 $\pm$ 0.0029	0.3154 $\pm$ 0.0206
default (d=160, L=5)	1.58M	0.0215 $\pm$ 0.0017	0.2899 $\pm$ 0.0095
shallow (L=2)	655k	0.0432 $\pm$ 0.0035	0.3198 $\pm$ 0.0224
deep (L=8)	2.51M	0.0158 $\pm$ 0.0009	0.3012 $\pm$ 0.0091
wide (d=320)	6.24M	0.0173 $\pm$ 0.0011	0.2956 $\pm$ 0.0146

303 The absolute family-transfer nMSE spans 0.2899–0.3198 across this $15\times$ parameter range — a spread
 304 comparable to each individual variant’s own 5-seed standard deviation (0.0091–0.0224), with no
 305 capacity-monotonic ordering (narrow < wide < deep < shallow in absolute family nMSE, tracking
 306 neither width nor depth). In-distribution reference nMSE, by contrast, *does* improve monotonically
 307 with capacity on both the width and depth axes, exactly as expected. Together this is evidence, not
 308 merely an absence of counter-evidence, that the cross-family gap is a genuine open problem rather
 309 than an artifact of the paper’s specific model size: scaling capacity roughly $4\times$ in either direction
 310 leaves it within noise of unchanged. A methodological note we flag explicitly: the $\times$ ref *ratio* is
 311 misleading as a cross-architecture comparison metric on its own — deep appears to have the worst
 312 relative gap ($19.1\times$) and shallow the best ($7.4\times$) largely because their reference nMSE differs by
 313 nearly $3\times$ (0.0158 vs. 0.0432), not because their absolute family-transfer error does (0.3012 vs.
 314 0.3198, well within either variant’s own std) — the same class of trap as the mean-vs-median artifact
 315 in Section 4.4; absolute nMSE, not the ratio, should be read as primary. This check does not cover
 316 context/query-length variation; a sweep over those axes is left to future work.

317 **On negative and inconvenient results.** Two of this paper’s four experiments (Sections 4.3 and 4.4)
 318 produced results that complicate rather than support the paper’s motivating narrative. We report both
 319 in full, including the Simpson’s-paradox confound our own first identifiability analysis missed and
 320 how the corrected analysis was obtained, because we believe an honestly-reported null result and a
 321 caught methodological error are more useful to the community building on this corpus than a quieter
 322 framing would be.

323 6 Broader Impact

324 PLANTFORGE is a synthetic corpus of control-plant dynamics; we are not aware of misuse vectors
 325 specific to this release beyond those generic to open ML tooling. Positive impact: broadening the
 326 training distribution available for in-context SysID research may reduce duplicated private data-
 327 generation effort across the field (the “demand proof” motivating this corpus, Section 2) and, via the
 328 released identifiability annotations, support excitation-design research relevant to real experimental
 329 efficiency in control engineering.

330 7 Conclusion

331 We released PLANTFORGE, a three-axis synthetic control-plant corpus with identifiability annota-
 332 tions, and subjected our own headline claim to progressively harder scrutiny: multi-seed error bars,
 333 zero-shot real-plant transfer on three real benchmarks, a classical baseline, a corrected statistical
 334 re-analysis of our identifiability hypothesis, and a 5-seed architecture ablation. The corpus’s central
 335 claim — broader training data narrows distribution-shift gaps for in-context SysID transformers —
 336 holds up under these checks, and the held-out-family gap specifically holds up as capacity-invariant
 337 across a $15\times$ parameter range, not an artifact of one model size. Two auxiliary claims we expected
 338 to hold did not: a classical baseline beats our trained transformers on every real-plant benchmark
 339 and several synthetic cells, and the identifiability annotations do not positively predict prediction

difficulty. One near-miss is worth naming directly: a single-seed preview of the Bouc-Wen real-plant result briefly suggested a third auxiliary reversal before the full 5-seed run corrected it back — exactly the failure mode our multi-seed standard exists to catch, and a useful reminder that “the model beat the baseline on this one run” is not the same claim as “on average.” We report all of this, with the analysis that uncovered each, as part of the release.

References

- Luis A. Aguirre, Bruno H. G. Barbosa, and Antônio P. Braga. Prediction and simulation errors in parameter estimation for nonlinear systems. *Mechanical Systems and Signal Processing*, 24(8): 2855–2867, 2010.
- Usman Akram and Haris Vikalo. Transformers as implicit state estimators: In-context learning in dynamical systems. *arXiv preprint arXiv:2410.16546*, 2024.
- Anonymous. PLANTFORGE corpus. link omitted for anonymous review; will be included in the camera-ready version, 2026.
- Logan Mondal Bhamidipaty, Tommy Bruzese, Caryn Tran, Rami Mrad, and Maxinder S. Kanwal. DynaDojo: An extensible platform for benchmarking scaling in dynamical system identification. In *Advances in Neural Information Processing Systems 36 (Datasets and Benchmarks Track)*, 2023.
- Stephen A. Billings. *Nonlinear System Identification: NARMAX Methods in the Time, Frequency, and Spatio-Temporal Domains*. Wiley, 2013. ISBN 9781119943594.
- Riccardo Busetto, Valentina Breschi, Marco Forgione, Dario Piga, and Simone Formentin. In-context learning of state estimators. In *20th IFAC Symposium on System Identification (SYSID)*, 2024. arXiv:2312.04509.
- Karl Cobbe, Christopher Hesse, Jacob Hilton, and John Schulman. Leveraging procedural generation to benchmark reinforcement learning. In *Proceedings of the 37th International Conference on Machine Learning (ICML)*, 2020. arXiv:1912.01588.
- Frank Cole, Yuxuan Zhao, Yulong Lu, and Tianhao Zhang. In-context learning of linear dynamical systems with transformers: Approximation bounds and depth-separation. *arXiv preprint arXiv:2502.08136*, 2025.
- Marco Forgione, Filippo Pura, and Dario Piga. From system models to class models: An in-context learning paradigm. *IEEE Control Systems Letters*, 2023.
- Timnit Gebru, Jamie Morgenstern, Briana Vecchione, Jennifer Wortman Vaughan, Hanna Wallach, Hal Daumé III, and Kate Crawford. Datasheets for datasets. *Communications of the ACM*, 64(12): 86–92, 2021. doi: 10.1145/3458723.
- Jayesh K. Gupta and Johannes Brandstetter. Towards multi-spatiotemporal-scale generalized PDE modeling. *arXiv preprint arXiv:2209.15616*, 2022.
- Håkan Hjalmarsson. From experiment design to closed-loop control. *Automatica*, 41(3):393–438, 2005.
- Pang Wei Koh, Shiori Sagawa, Henrik Marklund, Sang Michael Xie, Marvin Zhang, Akshay Balsubramani, Weihua Hu, Michihiro Yasunaga, Richard Lanus Phillips, Irena Gao, Tony Lee, Etienne David, Ian Stavness, Wei Guo, Berton A. Earnshaw, Imran S. Haque, Sara Beery, Jure Leskovec, Anshul Kundaje, Emma Pierson, Sergey Levine, Chelsea Finn, and Percy Liang. WILDS: A benchmark of in-the-wild distribution shifts. In *Proceedings of the 38th International Conference on Machine Learning (ICML)*, 2021.
- Lennart Ljung. *System Identification: Theory for the User*. Prentice Hall, 2nd edition, 1999. ISBN 0-13-656695-2.
- Fabio Muratore, Fabio Ramos, Greg Turk, Wenhao Yu, Michael Gienger, and Jan Peters. Robot learning from randomized simulations: A review. *Frontiers in Robotics and AI*, 9, 2022. doi: 10.3389/frobt.2022.799893.

388 Jean-Philippe Noël and Maarten Schoukens. Hysteretic benchmark with a dynamic nonlinearity, 2020.
389 Dataset, CC BY-SA 4.0. Benchmark design described in Noël & Schoukens (2016), Workshop on
390 Nonlinear System Identification Benchmarks, Brussels.

391 Dario Piga, Filippo Pura, and Marco Forgone. On the adaptation of in-context learners for system
392 identification. In *IFAC-PapersOnLine (Proceedings of the IFAC World Congress 2024)*, 2024.
393 arXiv:2312.04083.

394 Rik Pintelon and Johan Schoukens. *System Identification: A Frequency Domain Approach*. Wiley-
395 IEEE Press, 2nd edition, 2012. ISBN 9780470640371.

396 Matteo Rufolo, Dario Piga, and Marco Forgone. Distributionally robust minimization in meta-
397 learning for system identification. *arXiv preprint arXiv:2506.18074*, 2025a.

398 Matteo Rufolo, Dario Piga, Gabriele Maroni, and Marco Forgone. Enhanced transformer architecture
399 for in-context learning of dynamical systems. In *2025 European Control Conference (ECC)*, 2025b.
400 arXiv:2410.03291.

401 Johan Schoukens and Lennart Ljung. Nonlinear system identification: A user-oriented road map.
402 *IEEE Control Systems Magazine*, 39(6):28–99, 2019. doi: 10.1109/MCS.2019.2938121.

403 Jonas Sjöberg, Qinghua Zhang, Lennart Ljung, Albert Benveniste, Bernard Delyon, Pierre-
404 Yves Glorennec, Håkan Hjalmarsson, and Anatoli Juditsky. Nonlinear black-box modeling
405 in system identification: a unified overview. *Automatica*, 31(12):1691–1724, 1995. doi:
406 10.1016/0005-1098(95)00120-8.

407 Makoto Takamoto, Timothy Praditia, Raphael Leiteritz, Dan MacKinlay, Francesco Alesiani, Dirk
408 Pflüger, and Mathias Niepert. PDEBench: An extensive benchmark for scientific machine learning.
409 In *Advances in Neural Information Processing Systems 35 (Datasets and Benchmarks Track)*, 2022.
410 arXiv:2210.07182.

411 Josh Tobin, Rachel Fong, Alex Ray, Jonas Schneider, Wojciech Zaremba, and Pieter Abbeel. Domain
412 randomization for transferring deep neural networks from simulation to the real world. In *2017*
413 *IEEE/RSJ International Conference on Intelligent Robots and Systems (IROS)*, pages 23–30. IEEE,
414 2017.

415 Syed Izzat Ullah and Jose Baca. NanoBench: A multi-task benchmark dataset for nano-quadrotor
416 system identification, control, and state estimation. *arXiv preprint arXiv:2603.09908*, 2026.

417 Torbjörn Wigren and Johan Schoukens. Three free data sets for development and benchmarking in
418 nonlinear system identification. In *2013 European Control Conference (ECC)*, pages 2933–2938.
419 IEEE, 2013.

420 Martin Ziegler, Andres Felipe Posada-Moreno, Friedrich Solowjow, and Sebastian Trimpe. On foun-
421 dation models for dynamical systems from purely synthetic data. *arXiv preprint arXiv:2412.00395*,
422 2024.